

Wage LSTM Model - Training (1997-2020), Training (2021-2023)

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Input
from sklearn.model_selection import TimeSeriesSplit

import gdown
import time
```

```
In [ ]: # Load the dataset
url = 'https://drive.google.com/uc?id=1wjTiPLhi938Ro-jfjVHF0d_YPvsLaRc3'

# Download the file
output = 'data_wage.csv'
gdown.download(url, output, quiet=False)

# Check the file content
with open(output, 'r') as file:
    content = file.read()
    print("File content preview:")
    print(content[:500])

# Load the CSV file
try:
    data = pd.read_csv(output, delimiter=',')
    print(data.head())
except pd.errors.ParserError as e:
    print("Error parsing CSV file:", e)
```

```
Downloading...
From: https://drive.google.com/uc?id=1wjTiPLhi938Ro-jfjVHF0d_YPvsLaRc3
To: d:\OneDrive (Personal)\OneDrive\~ TMU 2023\CIND 820 - Big Data Analytics Project
\06 - Initial Results & Code (10%)\data_wage.csv
100%|██████████| 80.1M/80.1M [00:01<00:00, 43.0MB/s]
```

File content preview:

```
ref_date,geo,wages,type_of_work,sex,age_group,value,occupation_classification,noc,sex_binary,age_group_numeric,geo_code,date_ordinal,year,month
1997-01-01,newfoundland and labrador,average hourly wage rate,full-time employees,males,25 to 54 years,18.7,legislative and senior management occupations,00,1,1,210,729025,1997,1
```

```
1997-02-01,newfoundland and labrador,average hourly wage rate,full-time employees,males,25 to 54 years,18.48,legislative and senior management occupations,00,1,1,210,729056,1997
```

	ref_date	geo	wages	\
0	1997-01-01	newfoundland and labrador	average hourly wage rate	
1	1997-02-01	newfoundland and labrador	average hourly wage rate	
2	1997-03-01	newfoundland and labrador	average hourly wage rate	
3	1997-04-01	newfoundland and labrador	average hourly wage rate	
4	1997-05-01	newfoundland and labrador	average hourly wage rate	

	type_of_work	sex	age_group	value	\
0	full-time employees	males	25 to 54 years	18.70	
1	full-time employees	males	25 to 54 years	18.48	
2	full-time employees	males	25 to 54 years	27.87	
3	full-time employees	males	25 to 54 years	23.32	
4	full-time employees	males	25 to 54 years	23.08	

	occupation_classification	noc	sex_binary	\
0	legislative and senior management occupations	00	1	
1	legislative and senior management occupations	00	1	
2	legislative and senior management occupations	00	1	
3	legislative and senior management occupations	00	1	
4	legislative and senior management occupations	00	1	

	age_group_numeric	geo_code	date_ordinal	year	month
0	1	210	729025	1997	1
1	1	210	729056	1997	2
2	1	210	729084	1997	3
3	1	210	729115	1997	4
4	1	210	729145	1997	5

```
In [ ]: # Convert 'ref_date' to datetime
data['ref_date'] = pd.to_datetime(data['ref_date'], format='%Y-%m-%d')

# Label encode the 'occupation_classification' column
data['occupation_classification'] = data['occupation_classification'].astype(str)
label_encoder = LabelEncoder()
data['occupation_code'] = label_encoder.fit_transform(data['occupation_classification'])
```

```
In [ ]: # Normalize the data
scaler = MinMaxScaler(feature_range=(0, 1))
scaled_data = scaler.fit_transform(data[['value', 'date_ordinal', 'sex_binary', 'age_group_numeric', 'geo_code', 'date_ordinal', 'year', 'month']])

# Function to create sequences for forecasting
def create_sequences(data, seq_length):
    xs, ys = [], []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length, 0]
        xs.append(x)
```

```

        ys.append(y)
    return np.array(xs), np.array(ys)

# Define sequence length
SEQ_LENGTH = 12

# Split the data into training and testing sets based on the date
train_data = data[data['ref_date'] < '2021-01-01']
test_data = data[data['ref_date'] >= '2021-01-01']

# Normalize training and testing data separately
scaled_train_data = scaler.fit_transform(train_data[['value', 'date_ordinal', 'year', 'month_ordinal']])
scaled_test_data = scaler.transform(test_data[['value', 'date_ordinal', 'year', 'month_ordinal']])

```

```

In [ ]: # Create sequences for training and testing
X_train, y_train = create_sequences(scaled_train_data, SEQ_LENGTH)
X_test, y_test = create_sequences(scaled_test_data, SEQ_LENGTH)

# Define cross-validation procedure
tscv = TimeSeriesSplit(n_splits=5)
cv_mse_scores = []

# Perform cross-validation
for train_index, val_index in tscv.split(X_train):
    X_train_cv, X_val_cv = X_train[train_index], X_train[val_index]
    y_train_cv, y_val_cv = y_train[train_index], y_train[val_index]

    # Define the LSTM model for forecasting
    model = Sequential()
    model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
    model.add(LSTM(50, activation='relu'))
    model.add(Dense(1))
    model.compile(optimizer='adam', loss='mse')

    # Train the model
    start = time.time()
    history = model.fit(X_train_cv, y_train_cv, epochs=200, batch_size=32, validation_data=(X_val_cv, y_val_cv))
    end = time.time()

    # Evaluate the model
    val_loss = model.evaluate(X_val_cv, y_val_cv, verbose=0)
    cv_mse_scores.append(val_loss)

    print(f"Validation Loss: {val_loss}, Training time: {end - start} seconds")

print('Cross-Validation Mean Squared Error:', np.mean(cv_mse_scores))

```

```

Validation Loss: 0.0041073705069720745, Training time: 1485.4001214504242 seconds
Validation Loss: 0.0027351335156708956, Training time: 2564.582067012787 seconds
Validation Loss: 0.0022012367844581604, Training time: 3694.7265124320984 seconds
Validation Loss: 0.0035521364770829678, Training time: 5118.316879749298 seconds

```

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[7], line 23
    21 # Train the model
    22 start = time.time()
--> 23 history = model.fit(X_train_cv, y_train_cv, epochs=200, batch_size=32, validation_data=(X_val_cv, y_val_cv), verbose=0)
    24 end = time.time()
    26 # Evaluate the model

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\utils\traceback_utils.py:117, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    115 filtered_tb = None
    116 try:
--> 117     return fn(*args, **kwargs)
    118 except Exception as e:
    119     filtered_tb = _process_traceback_frames(e.__traceback__)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\backend\tensorflow\trainer.py:339, in TensorFlowTrainer.fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, validation_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, validation_steps, validation_batch_size, validation_freq)
    328 if getattr(self, "_eval_epoch_iterator", None) is None:
    329     self._eval_epoch_iterator = TFEPOCHIterator(
    330         x=val_x,
    331         y=val_y,
    (... )
    337         shuffle=False,
    338     )
--> 339 val_logs = self.evaluate(
    340     x=val_x,
    341     y=val_y,
    342     sample_weight=val_sample_weight,
    343     batch_size=validation_batch_size or batch_size,
    344     steps=validation_steps,
    345     callbacks=callbacks,
    346     return_dict=True,
    347     _use_cached_eval_dataset=True,
    348 )
    349 val_logs = {
    350     "val_" + name: val for name, val in val_logs.items()
    351 }
    352 epoch_logs.update(val_logs)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\utils\traceback_utils.py:117, in filter_traceback.<locals>.error_handler(*args, **kwargs)
    115 filtered_tb = None
    116 try:
--> 117     return fn(*args, **kwargs)
    118 except Exception as e:
    119     filtered_tb = _process_traceback_frames(e.__traceback__)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras

```

```

\src\backend\tensorflow\trainer.py:425, in TensorFlowTrainer.evaluate(self, x, y, batch_size, verbose, sample_weight, steps, callbacks, return_dict, **kwargs)
    423 for step, iterator in epoch_iterator.enumerate_epoch():
    424     callbacks.on_test_batch_begin(step)
--> 425     logs = self.test_function(iterator)
    426     logs = self._pythonify_logs(logs)
    427     callbacks.on_test_batch_end(step, logs)

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\util\traceback_utils.py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)

```

    148 filtered_tb = None
    149 try:
--> 150     return fn(*args, **kwargs)
    151 except Exception as e:
    152     filtered_tb = _process_traceback_frames(e.__traceback__)

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\polymorphic_function.py:833, in Function._call(self, *args, **kws)

```

    830 compiler = "xla" if self._jit_compile else "nonXla"
    832 with OptionalXlaContext(self._jit_compile):
--> 833     result = self._call(*args, **kws)
    835 new_tracing_count = self.experimental_get_tracing_count()
    836 without_tracing = (tracing_count == new_tracing_count)

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\polymorphic_function.py:878, in Function._call(self, *args, **kws)

```

    875 self._lock.release()
    876 # In this case we have not created variables on the first call. So we can
    877 # run the first trace but we should fail if variables are created.
--> 878 results = tracing_compilation.call_function(
    879     args, kws, self._variable_creation_config
    880 )
    881 if self._created_variables:
    882     raise ValueError("Creating variables on a non-first call to a function"
    883                      " decorated with tf.function.")

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\tracing_compilation.py:139, in call_function(args, kwargs, tracing_options)

```

    137 bound_args = function.function_type.bind(*args, **kwargs)
    138 flat_inputs = function.function_type.unpack_inputs(bound_args)
--> 139 return function._call_flat( # pylint: disable=protected-access
    140     flat_inputs, captured_inputs=function.captured_inputs
    141 )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\concrete_function.py:1322, in ConcreteFunction._call_flat(self, tensor_inputs, captured_inputs)

```

    1318 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
    1319 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE
    1320     and executing_eagerly):
    1321     # No tape is watching; skip to running the function.
-> 1322     return self._inference_function.call_preflattened(args)

```

```

1323 forward_backward = self._select_forward_and_backward_functions(
1324     args,
1325     possible_gradient_type,
1326     executing_eagerly)
1327 forward_function, args_with_tangents = forward_backward.forward()

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function.py:216, in AtomicFunction.call_preflattened(self, args)

```

214 def call_preflattened(self, args: Sequence[core.Tensor]) -> Any:
215     """Calls with flattened tensor inputs and returns the structured output.
--> 216     flat_outputs = self.call_flat(*args)
217     return self.function_type.pack_output(flat_outputs)

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function.py:251, in AtomicFunction.call_flat(self, *args)

```

249 with record.stop_recording():
250     if self._bound_context.executing_eagerly():
--> 251         outputs = self._bound_context.call_function(
252             self.name,
253             list(args),
254             len(self.function_type.flat_outputs),
255         )
256     else:
257         outputs = make_call_op_in_graph(
258             self,
259             list(args),
260             self._bound_context.function_call_options.as_attrs(),
261         )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\context.py:1500, in Context.call_function(self, name, tensor_inputs, num_outputs)

```

1498 cancellation_context = cancellation.context()
1499 if cancellation_context is None:
-> 1500     outputs = execute.execute(
1501         name.decode("utf-8"),
1502         num_outputs=num_outputs,
1503         inputs=tensor_inputs,
1504         attrs=attrs,
1505         ctx=self,
1506     )
1507 else:
1508     outputs = execute.execute_with_cancellation(
1509         name.decode("utf-8"),
1510         num_outputs=num_outputs,
1511         (...)
1514         cancellation_manager=cancellation_context,
1515     )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\execute.py:53, in quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)

```

51 try:

```

```

52     ctx.ensure_initialized()
---> 53     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
54                                             inputs, attrs, num_outputs)
55 except core._NotOkStatusException as e:
56     if name is not None:

```

KeyboardInterrupt:

```

In [ ]: # Create sequences for training and testing
X_train, y_train = create_sequences(scaled_train_data, SEQ_LENGTH)
X_test, y_test = create_sequences(scaled_test_data, SEQ_LENGTH)

```

```

In [ ]: # Define the LSTM model for forecasting
model = Sequential()
model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
model.add(LSTM(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')

print(model.summary())

# Train the model
start = time.time()
history = model.fit(X_train, y_train, epochs=500, batch_size=32, validation_split=0)
end = time.time()

# Convert elapsed time to minutes and seconds
elapsed_time = end - start
minutes = int(elapsed_time // 60)
seconds = int(elapsed_time % 60)

print(f"\nTraining time: {minutes} minutes and {seconds} seconds")

```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 50)	11,800
dense (Dense)	(None, 1)	51

Total params: 11,851 (46.29 KB)

Trainable params: 11,851 (46.29 KB)

Non-trainable params: 0 (0.00 B)

None

Epoch 1/500

1517/9427 ————— 27s 3ms/step - loss: 0.0096

```

-----
KeyboardInterrupt                                Traceback (most recent call last)
Cell In[5], line 16
    14 # Train the model
    15 start = time.time()
--> 16 history = model.fit(X_train, y_train, epochs=500, batch_size=32, validation_
split=0.2)
    17 end = time.time()
    19 # Convert elapsed time to minutes and seconds

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras
\src\utils\traceback_utils.py:117, in filter_traceback.<locals>.error_handler(*args,
**kwargs)
    115 filtered_tb = None
    116 try:
--> 117     return fn(*args, **kwargs)
    118 except Exception as e:
    119     filtered_tb = _process_traceback_frames(e.__traceback__)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras
\src\backend\tensorflow\trainer.py:314, in TensorFlowTrainer.fit(self, x, y, batch_s
ize, epochs, verbose, callbacks, validation_split, validation_data, shuffle, class_w
eight, sample_weight, initial_epoch, steps_per_epoch, validation_steps, validation_b
atch_size, validation_freq)
    312 for step, iterator in epoch_iterator.enumerate_epoch():
    313     callbacks.on_train_batch_begin(step)
--> 314     logs = self.train_function(iterator)
    315     logs = self._pythonify_logs(logs)
    316     callbacks.on_train_batch_end(step, logs)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensor
flow\python\util\traceback_utils.py:150, in filter_traceback.<locals>.error_handler
(*args, **kwargs)
    148 filtered_tb = None
    149 try:
--> 150     return fn(*args, **kwargs)
    151 except Exception as e:
    152     filtered_tb = _process_traceback_frames(e.__traceback__)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensor
flow\python\eager\polymorphic_function\polymorphic_function.py:833, in Function._ca
ll__(self, *args, **kws)
    830 compiler = "xla" if self._jit_compile else "nonXla"
    832 with OptionalXlaContext(self._jit_compile):
--> 833     result = self._call(*args, **kws)
    835 new_tracing_count = self.experimental_get_tracing_count()
    836 without_tracing = (tracing_count == new_tracing_count)

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensor
flow\python\eager\polymorphic_function\polymorphic_function.py:878, in Function._cal
l(self, *args, **kws)
    875 self._lock.release()
    876 # In this case we have not created variables on the first call. So we can
    877 # run the first trace but we should fail if variables are created.
--> 878 results = tracing_compilation.call_function(
    879     args, kws, self._variable_creation_config

```



```

880 )
881 if self._created_variables:
882     raise ValueError("Creating variables on a non-first call to a function"
883                       " decorated with tf.function.")

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\tracing_compilation.py:139, in call_function(args, kwargs, tracing_options)

```

137 bound_args = function.function_type.bind(*args, **kwargs)
138 flat_inputs = function.function_type.unpack_inputs(bound_args)
--> 139 return function._call_flat( # pylint: disable=protected-access
140     flat_inputs, captured_inputs=function.captured_inputs
141 )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\concrete_function.py:1322, in ConcreteFunction._call_flat(self, tensor_inputs, captured_inputs)

```

1318 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
1319 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE
1320     and executing_eagerly):
1321     # No tape is watching; skip to running the function.
-> 1322     return self._inference_function.call_prelattened(args)
1323 forward_backward = self._select_forward_and_backward_functions(
1324     args,
1325     possible_gradient_type,
1326     executing_eagerly)
1327 forward_function, args_with_tangents = forward_backward.forward()

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function.py:216, in AtomicFunction.call_prelattened(self, args)

```

214 def call_prelattened(self, args: Sequence[core.Tensor]) -> Any:
215     """Calls with flattened tensor inputs and returns the structured output.
216     """
--> 216     flat_outputs = self.call_flat(*args)
217     return self.function_type.pack_output(flat_outputs)

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\polymorphic_function\atomic_function.py:251, in AtomicFunction.call_flat(self, *args)

```

249 with record.stop_recording():
250     if self._bound_context.executing_eagerly():
--> 251         outputs = self._bound_context.call_function(
252             self.name,
253             list(args),
254             len(self.function_type.flat_outputs),
255         )
256     else:
257         outputs = make_call_op_in_graph(
258             self,
259             list(args),
260             self._bound_context.function_call_options.as_attrs(),
261         )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\context.py:1500, in Context.call_function(self, name, tensor_input

```

s, num_outputs)
1498 cancellation_context = cancellation.context()
1499 if cancellation_context is None:
-> 1500     outputs = execute.execute(
1501         name.decode("utf-8"),
1502         num_outputs=num_outputs,
1503         inputs=inputs,
1504         attrs=attrs,
1505         ctx=self,
1506     )
1507 else:
1508     outputs = execute.execute_with_cancellation(
1509         name.decode("utf-8"),
1510         num_outputs=num_outputs,
1511         (...)
1512     )
1513     cancellation_manager=cancellation_context,
1514 )
1515 )

```

File c:\Users\nesha\AppData\Local\Programs\Python\Python311\Lib\site-packages\tensorflow\python\eager\execute.py:53, in quick_execute(op_name, num_outputs, inputs, attrs, s, ctx, name)

```

51 try:
52     ctx.ensure_initialized()
---> 53     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
54         inputs, attrs, num_outputs)
55 except core._NotOkStatusException as e:
56     if name is not None:

```

KeyboardInterrupt:

```

In [ ]: # Define the smoothing function
def smooth_curve(points, factor=0.9):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

# Retrieve loss and validation loss from history
loss = history.history['loss']
val_loss = history.history['val_loss']

# Smooth the loss curves
smoothed_loss = smooth_curve(loss)
smoothed_val_loss = smooth_curve(val_loss)

# Plot smoothed training and validation loss
plt.plot(range(1, len(smoothed_loss) + 1), smoothed_loss, label='Smoothed Training')
plt.plot(range(1, len(smoothed_val_loss) + 1), smoothed_val_loss, label='Smoothed V')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

```

```

plt.show()

# Identify the epoch with the lowest validation loss
best_epoch = np.argmin(smoothed_val_loss) + 1
print(f"Best epoch based on validation loss: {best_epoch}")


# Refine the LSTM model
model = Sequential()
model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
model.add(LSTM(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')

print(model.summary())

# Train the model
start = time.time()
history = model.fit(X_train, y_train, epochs=best_epoch, batch_size=32, validation_
end = time.time())

print(f"\nTraining time: {minutes} minutes and {seconds} seconds")

```



Best epoch based on validation loss: 56
 Model: "sequential_1"

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 50)	11,800
dense_1 (Dense)	(None, 1)	51


Total params: 11,851 (46.29 KB)

Trainable params: 11,851 (46.29 KB)

Non-trainable params: 0 (0.00 B)

None

Epoch 1/56

9427/9427  31s 3ms/step - loss: 0.0047 - val_loss: 0.0031


Epoch 2/56

9427/9427  29s 3ms/step - loss: 0.0027 - val_loss: 0.0031


Epoch 3/56

9427/9427  29s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 4/56

9427/9427  29s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 5/56

9427/9427  29s 3ms/step - loss: 0.0026 - val_loss: 0.0029

Epoch 6/56

9427/9427  29s 3ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 7/56

9427/9427  29s 3ms/step - loss: 0.0025 - val_loss: 0.0029


Epoch 8/56

9427/9427  29s 3ms/step - loss: 0.0025 - val_loss: 0.0029

Epoch 9/56

9427/9427  30s 3ms/step - loss: 0.0024 - val_loss: 0.0031


Epoch 10/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0029


Epoch 11/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0030


Epoch 12/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0028


Epoch 13/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0029


Epoch 14/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0030


Epoch 15/56

9427/9427  30s 3ms/step - loss: 0.0024 - val_loss: 0.0030


Epoch 16/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0028

Epoch 17/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0028

Epoch 18/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0031


Epoch 19/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0029


Epoch 20/56

9427/9427  30s 3ms/step - loss: 0.0024 - val_loss: 0.0028

Epoch 21/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0028

Epoch 22/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0028


Epoch 23/56

9427/9427  29s 3ms/step - loss: 0.0024 - val_loss: 0.0030


Epoch 24/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0029

Epoch 25/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0028

Epoch 26/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0030

Epoch 27/56

9427/9427  29s 3ms/step - loss: 0.0023 - val_loss: 0.0030

Epoch 28/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 29/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 30/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 31/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 32/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 33/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 34/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 35/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 36/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 37/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 38/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 39/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 40/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 41/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 42/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 43/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 44/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 45/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0029
Epoch 46/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 47/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 48/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 49/56

9427/9427 ————— 29s 3ms/step - loss: 0.0022 - val_loss: 0.0027
Epoch 50/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 51/56

9427/9427 ————— 29s 3ms/step - loss: 0.0023 - val_loss: 0.0027
Epoch 52/56

9427/9427 ————— 30s 3ms/step - loss: 0.0023 - val_loss: 0.0028
Epoch 53/56

9427/9427 ————— 29s 3ms/step - loss: 21.4818 - val_loss: 0.0064
Epoch 54/56

9427/9427 ————— 30s 3ms/step - loss: 0.0112 - val_loss: 0.0036
Epoch 55/56

9427/9427 ————— 30s 3ms/step - loss: 0.0028 - val_loss: 0.0029
Epoch 56/56

9427/9427 ————— 30s 3ms/step - loss: 0.0024 - val_loss: 0.0028

Training time: 256 minutes and 14 seconds

```
In [ ]: # Evaluate the model
loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {loss}')

# Make predictions
predictions = model.predict(X_test)

# Inverse transform the predictions and the actual values
predictions_inv = scaler.inverse_transform(np.concatenate((predictions, np.zeros((p
y_test_inv = scaler.inverse_transform(np.concatenate((y_test.reshape(-1, 1), np.zer

# Compare the first few predictions with the actual values
comparison = pd.DataFrame({'Actual': y_test_inv, 'Predicted': predictions_inv})
print(comparison.head())
```

1516/1516 ————— 2s 1ms/step - loss: 0.0225

Test Loss: 0.022555621340870857

1516/1516 ————— 2s 1ms/step

	Actual	Predicted
0	48.51	21.124145
1	49.43	31.216181
2	49.78	33.773115
3	50.11	47.519050
4	52.02	41.261429

```
In [ ]: # Wage LSTM Model - Training (1997-2020), Testing (2021-2023)
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from sklearn.model_selection import TimeSeriesSplit
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Input
import gdown
import time

# Load the dataset
url = 'https://drive.google.com/uc?id=1wjTiPLhi938Ro-jfjVHF0d_YPvsLaRc3'

# Download the file
output = 'data_wage.csv'
gdown.download(url, output, quiet=False)

# Check the file content
with open(output, 'r') as file:
    content = file.read()
    print("File content preview:")
    print(content[:500])

# Load the CSV file
```

```

try:
    data_wage = pd.read_csv(output, delimiter=',')
    print(data_wage.head())
except pd.errors.ParserError as e:
    print("Error parsing CSV file:", e)

# Convert 'ref_date' to datetime
data_wage['ref_date'] = pd.to_datetime(data_wage['ref_date'], format='%Y-%m-%d')

# Change 'ref_date' format to '%Y-%m'
data_wage['ref_date'] = data_wage['ref_date'].dt.to_period('M').dt.to_timestamp()

# Label encode the 'occupation_classification' column
data_wage['occupation_classification'] = data_wage['occupation_classification'].astype(str)
label_encoder = LabelEncoder()
data_wage['occupation_code'] = label_encoder.fit_transform(data_wage['occupation_classification'])

# Normalize the data
scaler = MinMaxScaler(feature_range=(0, 1))
features = ['value', 'date_ordinal', 'sex_binary', 'age_group_numeric', 'geo_code', 'occupation_code']
scaled_data = scaler.fit_transform(data_wage[features])

# Function to create sequences for forecasting
def create_sequences(data, seq_length):
    xs, ys = [], []
    for i in range(len(data) - seq_length):
        x = data[i:i+seq_length]
        y = data[i+seq_length, 0]
        xs.append(x)
        ys.append(y)
    return np.array(xs), np.array(ys)

# Define sequence length
SEQ_LENGTH = 12

# Split the data into training and testing sets based on the date
train_data = data_wage[data_wage['ref_date'] < '2021-01']
test_data = data_wage[data_wage['ref_date'] >= '2021-01']

# Normalize training and testing data separately
scaled_train_data = scaler.fit_transform(train_data[features])
scaled_test_data = scaler.transform(test_data[features])

# Create sequences for training and testing
X_train, y_train = create_sequences(scaled_train_data, SEQ_LENGTH)
X_test, y_test = create_sequences(scaled_test_data, SEQ_LENGTH)

# Define cross-validation procedure
tscv = TimeSeriesSplit(n_splits=5)
cv_mse_scores = []

# Perform cross-validation
for train_index, val_index in tscv.split(X_train):
    X_train_cv, X_val_cv = X_train[train_index], X_train[val_index]
    y_train_cv, y_val_cv = y_train[train_index], y_train[val_index]

```



```

# Define the LSTM model for forecasting
model = Sequential()
model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
model.add(LSTM(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')

# Train the model
start = time.time()
history = model.fit(X_train_cv, y_train_cv, epochs=100, batch_size=32, validation_split=0.1)
end = time.time()

# Evaluate the model
val_loss = model.evaluate(X_val_cv, y_val_cv, verbose=0)
cv_mse_scores.append(val_loss)

print(f"Validation Loss: {val_loss}, Training time: {end - start} seconds")

print('Cross-Validation Mean Squared Error:', np.mean(cv_mse_scores))

# Train the final model on the entire training set
model = Sequential()
model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
model.add(LSTM(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')

print(model.summary())

start = time.time()
history = model.fit(X_train, y_train, epochs=100, batch_size=32, validation_split=0.1)
end = time.time()

# Convert elapsed time to minutes and seconds
elapsed_time = end - start
minutes = int(elapsed_time // 60)
seconds = int(elapsed_time % 60)

print(f"\nTraining time: {minutes} minutes and {seconds} seconds")

# Retrieve loss and validation loss from history
loss = history.history['loss']
val_loss = history.history['val_loss']

# Define the smoothing function
def smooth_curve(points, factor=0.9):
    smoothed_points = []
    for point in points:
        if smoothed_points:
            previous = smoothed_points[-1]
            smoothed_points.append(previous * factor + point * (1 - factor))
        else:
            smoothed_points.append(point)
    return smoothed_points

# Smooth the Loss curves

```

```

smoothed_loss = smooth_curve(loss)
smoothed_val_loss = smooth_curve(val_loss)

# Plot smoothed training and validation loss
plt.plot(range(1, len(smoothed_loss) + 1), smoothed_loss, label='Smoothed Training')
plt.plot(range(1, len(smoothed_val_loss) + 1), smoothed_val_loss, label='Smoothed Validation')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

# Identify the epoch with the lowest validation loss
best_epoch = np.argmin(smoothed_val_loss) + 1
print(f"Best epoch based on validation loss: {best_epoch}")

# Refine the LSTM model
model = Sequential()
model.add(Input(shape=(SEQ_LENGTH, X_train.shape[2])))
model.add(LSTM(50, activation='relu'))
model.add(Dense(1))
model.compile(optimizer='adam', loss='mse')

print(model.summary())

# Train the model
start = time.time()
history = model.fit(X_train, y_train, epochs=best_epoch, batch_size=32, validation_data=(X_test, y_test))
end = time.time()

print(f"\nTraining time: {minutes} minutes and {seconds} seconds")

# Evaluate the model
loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {loss}')

# Make predictions
predictions = model.predict(X_test)

# Inverse transform the predictions and the actual values
predictions_inv = scaler.inverse_transform(np.concatenate((predictions, np.zeros((predictions.shape[0], 1)))))
y_test_inv = scaler.inverse_transform(np.concatenate((y_test.reshape(-1, 1), np.zeros((predictions.shape[0], 1)))))

# Compare the first few predictions with the actual values
comparison = pd.DataFrame({'Actual': y_test_inv, 'Predicted': predictions_inv})
print(comparison.head())

```

Downloading...

From: https://drive.google.com/uc?id=1wjTiPLhi938Ro-jfjVHF0d_YPvsLaRc3

To: d:\OneDrive (Personal)\OneDrive\~ TMU 2023\CIND 820 - Big Data Analytics Project\06 - Initial Results & Code (10%)\data_wage.csv

100%|██████████| 80.1M/80.1M [00:02<00:00, 38.7MB/s]

File content preview:

```
ref_date,geo,wages,type_of_work,sex,age_group,value,occupation_classification,noc,sex_binary,age_group_numeric,geo_code,date_ordinal,year,month
1997-01-01,newfoundland and labrador,average hourly wage rate,full-time employees,males,25 to 54 years,18.7,legislative and senior management occupations,00,1,1,210,729025,1997,1
```

```
1997-02-01,newfoundland and labrador,average hourly wage rate,full-time employees,males,25 to 54 years,18.48,legislative and senior management occupations,00,1,1,210,729056,1997
```

```

      ref_date      geo      wages \
0  1997-01-01  newfoundland and labrador  average hourly wage rate
1  1997-02-01  newfoundland and labrador  average hourly wage rate
2  1997-03-01  newfoundland and labrador  average hourly wage rate
3  1997-04-01  newfoundland and labrador  average hourly wage rate
4  1997-05-01  newfoundland and labrador  average hourly wage rate

```

```

      type_of_work  sex  age_group  value \
0  full-time employees  males  25 to 54 years  18.70
1  full-time employees  males  25 to 54 years  18.48
2  full-time employees  males  25 to 54 years  27.87
3  full-time employees  males  25 to 54 years  23.32
4  full-time employees  males  25 to 54 years  23.08

```

```

      occupation_classification  noc  sex_binary \
0  legislative and senior management occupations  00      1
1  legislative and senior management occupations  00      1
2  legislative and senior management occupations  00      1
3  legislative and senior management occupations  00      1
4  legislative and senior management occupations  00      1

```

```

      age_group_numeric  geo_code  date_ordinal  year  month
0          1          210          729025  1997      1
1          1          210          729056  1997      2
2          1          210          729084  1997      3
3          1          210          729115  1997      4
4          1          210          729145  1997      5

```

Validation Loss: 0.0027075393591076136, Training time: 710.9991371631622 seconds

Validation Loss: 0.0027307418640702963, Training time: 1157.2986750602722 seconds

Validation Loss: 0.0023171829525381327, Training time: 1626.4196791648865 seconds

Validation Loss: 0.003146085189655423, Training time: 2079.2018325328827 seconds

Validation Loss: 0.2398611158132553, Training time: 2571.675872325897 seconds

Cross-Validation Mean Squared Error: 0.05015253303572535

Model: "sequential_5"

Layer (type)	Output Shape	Param #
lstm_5 (LSTM)	(None, 50)	11,400
dense_5 (Dense)	(None, 1)	51


Total params: 11,451 (44.73 KB)

Trainable params: 11,451 (44.73 KB)

Non-trainable params: 0 (0.00 B)

None

Epoch 1/100

9427/9427  28s 3ms/step - loss: 0.0049 - val_loss: 0.0031


Epoch 2/100

9427/9427  27s 3ms/step - loss: 0.0027 - val_loss: 0.0031


Epoch 3/100

9427/9427  26s 3ms/step - loss: 0.0027 - val_loss: 0.0030


Epoch 4/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 5/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 6/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 7/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 8/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 9/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 10/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0032


Epoch 11/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 12/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 13/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 14/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 15/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 16/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 17/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 18/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 19/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 20/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 21/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 22/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 23/100

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0029


Epoch 24/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 25/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0029

Epoch 26/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 27/100

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 28/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 29/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0031
Epoch 30/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 31/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0031
Epoch 32/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 33/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 34/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 35/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 36/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 37/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 38/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 39/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 40/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 41/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 42/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 43/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 44/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 45/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 46/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 47/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 48/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0031
Epoch 49/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 50/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 51/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 52/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 53/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 54/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 55/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 56/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 57/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 58/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 59/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 60/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 61/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 62/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 63/100

9427/9427 ————— 27s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 64/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 65/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 66/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 67/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 68/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 69/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 70/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 71/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 72/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 73/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 74/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 75/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 76/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 77/100

9427/9427 ————— 26s 3ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 78/100

9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 79/100


9427/9427 ————— 27s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 80/100


9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 81/100


9427/9427 ————— 26s 3ms/step - loss: 0.0033 - val_loss: 0.0029
Epoch 82/100


9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 83/100


9427/9427 ————— 26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 84/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 85/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 86/100


9427/9427  27s 3ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 87/100


9427/9427  27s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 88/100


9427/9427  27s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 89/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 90/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 91/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 92/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 93/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0033
Epoch 94/100


9427/9427  27s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 95/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 96/100

9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 97/100


9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 98/100

9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 99/100

9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 100/100

9427/9427  26s 3ms/step - loss: 0.0024 - val_loss: 0.0029

Training time: 43 minutes and 39 seconds



Best epoch based on validation loss: 89
Model: "sequential_6"

Layer (type)	Output Shape	Param #
lstm_6 (LSTM)	(None, 50)	11,400
dense_6 (Dense)	(None, 1)	51

Total params: 11,451 (44.73 KB)
Trainable params: 11,451 (44.73 KB)
Non-trainable params: 0 (0.00 B)

None

Epoch 1/89

9427/9427  29s 3ms/step - loss: 0.0058 - val_loss: 0.0031


Epoch 2/89

9427/9427  26s 3ms/step - loss: 0.0027 - val_loss: 0.0030


Epoch 3/89

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 4/89

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0031

Epoch 5/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0031

Epoch 6/89

9427/9427  26s 3ms/step - loss: 0.0026 - val_loss: 0.0031


Epoch 7/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 8/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 9/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 10/89

9427/9427  26s 3ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 11/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 12/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030


Epoch 13/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0031


Epoch 14/89

9427/9427  27s 3ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 15/89

9427/9427  27s 3ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 16/89

9427/9427  32s 3ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 17/89

9427/9427  7331s 778ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 18/89

9427/9427  63s 7ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 19/89

9427/9427  57s 6ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 20/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 21/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 22/89

9427/9427  59s 6ms/step - loss: 0.0026 - val_loss: 0.0030

Epoch 23/89

9427/9427  57s 6ms/step - loss: 0.0025 - val_loss: 0.0030


Epoch 24/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 25/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 26/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0029

Epoch 27/89

9427/9427  58s 6ms/step - loss: 0.0025 - val_loss: 0.0030

Epoch 28/89

```

9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 29/89
9427/9427 ██████████ 57s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 30/89
9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 31/89
9427/9427 ██████████ 57s 6ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 32/89
9427/9427 ██████████ 59s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 33/89
9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 34/89
9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 35/89
9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 36/89
9427/9427 ██████████ 58s 6ms/step - loss: 0.0025 - val_loss: 0.0031
Epoch 37/89
9427/9427 ██████████ 60s 6ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 38/89
9427/9427 ██████████ 7346s 779ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 39/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 40/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 41/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 42/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 43/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 44/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 45/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 46/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 47/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 48/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 49/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 50/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 51/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 52/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 53/89
9427/9427 ██████████ 7336s 778ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 54/89
9427/9427 ██████████ 89s 9ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 55/89
9427/9427 ██████████ 75s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 56/89

```

```

9427/9427 ██████████ 75s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 57/89
9427/9427 ██████████ 75s 8ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 58/89
9427/9427 ██████████ 75s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 59/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 60/89
9427/9427 ██████████ 76s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 61/89
9427/9427 ██████████ 75s 8ms/step - loss: 0.0025 - val_loss: 0.0031
Epoch 62/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 63/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 64/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 65/89
9427/9427 ██████████ 75s 8ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 66/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0030
Epoch 67/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 68/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 69/89
9427/9427 ██████████ 74s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 70/89
9427/9427 ██████████ 7342s 779ms/step - loss: 0.0025 - val_loss: 0.0030
Epoch 71/89
9427/9427 ██████████ 77s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 72/89
9427/9427 ██████████ 73s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 73/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 74/89
9427/9427 ██████████ 73s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 75/89
9427/9427 ██████████ 73s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 76/89
9427/9427 ██████████ 73s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 77/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 78/89
9427/9427 ██████████ 73s 8ms/step - loss: 0.0024 - val_loss: 0.0029
Epoch 79/89
9427/9427 ██████████ 72s 8ms/step - loss: 141.0687 - val_loss: 0.0283
Epoch 80/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.1419 - val_loss: 0.0062
Epoch 81/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.0101 - val_loss: 0.0043
Epoch 82/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.0033 - val_loss: 0.0031
Epoch 83/89
9427/9427 ██████████ 72s 8ms/step - loss: 0.0027 - val_loss: 0.0031
Epoch 84/89

```

9427/9427  **72s** 8ms/step - loss: 0.0026 - val_loss: 0.0029
Epoch 85/89

9427/9427  **72s** 8ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 86/89

9427/9427  **4008s** 425ms/step - loss: 0.0025 - val_loss: 0.0029
Epoch 87/89

9427/9427  **220424s** 23s/step - loss: 5.6695 - val_loss: 2.9401
Epoch 88/89

9427/9427  **34s** 4ms/step - loss: 0.1000 - val_loss: 0.0212
Epoch 89/89

9427/9427  **41s** 4ms/step - loss: 0.0072 - val_loss: 0.0892

Training time: 43 minutes and 39 seconds

1516/1516  **2s** 1ms/step - loss: 0.2851

Test Loss: 0.852683961391449

1516/1516  **3s** 2ms/step

	Actual	Predicted
0	48.51	51.434662
1	49.43	48.974869
2	49.78	50.224376
3	50.11	48.432188
4	52.02	47.580685